
PrimaLearn in Neural Program Synthesis: Spelke-Initialized Library Learning Discovers Cross-Domain Abstractions

Purna Sainath Reddy V

*Department of Computer Science and Engineering (Data Science),
Manipal Institute of Technology Bengaluru, Manipal Academy of Higher Education*

vpurnareddy6@gmail.com

ABSTRACT

We present *bootstrap-substrate*, the first end-to-end computational instantiation of Carey's PrimaLearn theory within a neural program-synthesis architecture. The system initializes a typed domain-specific language (DSL) with primitives from Spelke's core-knowledge systems (OBJECTS, FORMS, NUMBER) and runs a DreamCoder-style wake-sleep loop with Stitch anti-unification over the ARC-AGI benchmark. We demonstrate three properties required by the bootstrapping account: (1) the compressor discovers *novel cross-system abstractions* bridging FORMS and OBJECTS, specifically `abs0,2`, a parameterized combinator with MDL savings of 25, reused across 4 programs on real ARC tasks; (2) these abstractions *compound*: an OBJECTS abstraction (`abs0,10`) discovered in cycle 0 enables solving task `b9b7f026` in cycle 1 via composition with OBJECTS gravity operations, a program unreachable without the learned abstraction; and (3) the Spelke-initialized DSL *outperforms* both a matched-cardinality Generic DSL and the VIMRL objectness-only DSL on ARC-AGI: 46/400 tasks (11.5%) vs. 42/400 (10.5%) for both baselines (statistical significance assessed via McNemar's test, 5-seed multi-run analysis), with 14 unique abstractions discovered (including 1 cross-system) vs. 0 for each baseline. The library grows from 143 to 171 primitives across 3 cycles while baseline libraries remain static. All results derive from a single canonical experiment run with full reproducibility. A 5-seed multi-run analysis (McNemar's test) yields $\chi^2=1.3333$, $p=0.2482$ (Spelke vs. Generic): the +4 task advantage is consistent across all 5 seeds but is *not statistically significant* at conventional thresholds ($p<0.05$), reflecting a power limitation at the 400-task scale. A sample efficiency curve (`task_count` $\in$ {10, 25, 50, 100, 200, 400}) provides complementary monotonic evidence for the advantage.

1. Introduction

How does a mind acquire concepts that go beyond its starting vocabulary? Carey's theory of *PrimaLearn* [4] offers a precise answer: placeholder symbols from distinct conceptual systems are first used without full interpretation, then cross-system modelling processes gradually bind them together, producing new concepts with expressive power strictly greater than any source system. The canonical example is the child's construction of the integer concept by bootstrapping from the object-tracking system (which represents small exact cardinalities) and the approximate-number system (which represents large magnitudes), with the placeholder symbols *one*, *two*, *three* acting as the bridge [4, 5].

Spelke's core-knowledge framework [20] provides the vocabulary of starting systems: OBJECTS (cohesion, persistence, contact), FORMS (geometry, symmetry, periodicity),

NUMBER (object-tracking system, analog magnitude system), AGENTS, PERSONS, and PLACES. These are posited to be innate, domain-specific, and the raw material from which bootstrapping builds richer concepts.

Despite its theoretical prominence, PrimaLearn has never been given a mechanistic computational instantiation. Prior work in program induction [8, 15, 19] learns libraries of reusable combinators but does not ground them in Spelke systems, does not track cross-system provenance, and does not claim to demonstrate the Carey bootstrapping sequence. Work in developmental AI [1, 6, 14] invokes the framework rhetorically without implementing the placeholder-binding mechanism.

We fill this gap. Our contributions are:

1. **The first Spelke-typed DSL for visual program induction.**
143 primitives organized into the three active core-knowl-

edge systems (OBJECTS, FORMS, NUMBER) with explicit system provenance tracked on every primitive.

2. **Cross-system abstraction discovery.** A DreamCoder wake-sleep loop [8] with Stitch anti-unification [3] discovers cross-system abstractions bridging FORMS and OBJECTS, both on synthetic tasks (`abs0,0`, MDL savings 124) and on real ARC tasks (`abs0,2`, MDL savings 25). These are recovered by MDL compression, not heuristics.
3. **Empirical demonstration of compounding.** On the Spelke Bootstrap Suite v2 (§4.3), 75 tasks unsolvable in cycle 0 become solvable in cycle 1 via three cross-system abstractions discovered in cycle 0. On ARC, task `b9b7f026` is solved in cycle 1 via composition of a learned OBJECTS abstraction with OBJECTS gravity operations, a program unreachable within the enumerator’s cost budget without the learned abstraction.
4. **Three-way baseline comparison.** The Spelke DSL (143 prims) achieves 46/400 tasks (11.5%) on ARC-AGI versus 42/400 (10.5%) for a matched-cardinality Generic DSL and 42/400 (10.5%) for the VIMRL objectness-only DSL [1]. Spelke discovers 14 abstractions across 3 cycles (library 143→171); VIMRL discovers 0. The advantage is attributable to OBJECTS-system grounding enabling cross-system program composition.

2. Background and Related Work

2.1 PrimaLearn

Carey [4] defines PrimaLearn as a process in which placeholder symbols (terms used before their meaning is fully determined) are integrated across incommensurable representational systems through analogical and structural mapping processes. The resulting concept is not reducible to either source system: the integer concept is neither pure cardinality-tracking nor pure magnitude estimation, but a new structure that emerges from their coupling. Spelke and colleagues [20, 5] provide the empirical base: infant studies demonstrate that the six core-knowledge systems are present from birth, domain-specific, and computation-ready, making them the natural starting vocabulary for any bootstrapping account.

2.2 Program Synthesis and Library Learning

Ellis et al. [8] introduced DreamCoder, a wake-sleep algorithm for program induction: a *wake* phase solves tasks using a recognition network to guide enumeration; an *abstraction sleep* phase compresses the solved-program corpus into a

growing library; a *dreaming sleep* phase trains the recognition network on imagined programs. The library grows monotonically, and each cycle’s abstractions become first-class primitives in the next. Stitch [3] provides a scalable anti-unification algorithm that finds the MDL-optimal library given a corpus, replacing the slower compression step in the original DreamCoder.

The ARC-AGI benchmark [6] (Abstraction and Reasoning Corpus) presents 400 training tasks, each a small set of input-output grid pairs, designed to require core-knowledge reasoning without large-data pretraining. It has become the standard evaluation for human-like abstraction. Competitive approaches include hand-crafted DSLs [11], learned program synthesis [16], and large language models with tool use [21]; none are grounded in the developmental cognitive science literature.

2.3 Cognitive-AI Alignment

Ainooson et al. [1] present VIMRL, an objectness-centric DSL for ARC with 45 primitives, the closest prior work in spirit. They do not implement number or forms primitives, do not run a library-learning loop, and do not relate their results to Carey’s bootstrapping. Lake et al. [14] propose Building Machines That Learn and Think Like People but focus on model-based reinforcement learning. Our work is the first to translate the specific Carey-Spelke theoretical framework into working software with measurable predictions.

3. Method

3.1 The Spelke-Typed DSL

We define a simply-typed lambda calculus over ARC grids. Base types are `grid`, `object`, `color`, `int`, `bool`, `list[·]`, and `point`. Every primitive carries a *system tag* $\in \{\text{OBJECTS, FORMS, NUMBER, GLUE}\}$ inherited from its Spelke provenance. The full library contains 143 typed primitives:

System	Prims	Representative primitives
OBJECTS	38	<code>extract_objects</code> , <code>obj_largest</code> , <code>render_object</code> , <code>gravity×4</code> , <code>fill_interior</code>
FORMS	31	<code>rotate90/180/270</code> , <code>flip_h/v</code> , <code>transform</code> , <code>extract_tile</code> , <code>scale_up/down</code>
NUMBER	41	<code>count_objects</code> , <code>obj_size</code> , OTS cap-3, ANS Weber ratio

Table 1. Spelke-typed DSL primitive counts by system. GLUE includes generic higher-order combinators and the cross-domain analogy primitive.

Each primitive's type signature is enforced during enumeration: the type-directed enumerator [8] only constructs well-typed programs, which radically prunes the search space for cross-type compositions.

3.2 Wake-Sleep Architecture




Figure 1. Wake-sleep architecture. The Spelke-typed DSL (143 primitives across OBJECTS, FORMS, NUMBER, and GLUE) feeds the AST solver; Stitch discovers shared structure across solved programs; discovered abstractions become first-class primitives in subsequent cycles. The cross-system abstraction box (red) is the Carey signature: it emerges from the compression step, not from any prior over cross-system compositions.

Wake phase. For each task, the system first attempts 22 hand-coded AST-producing strategies organized by Spelke system (single-primitive FORMS transforms, OBJECTS pipelines, color operations, gravity, symmetry completion, scaling, tiling, and five cross-system strategies including FORMS-on-extracted-object and NUMBER-recolor). Strategies return typed program ASTs (not opaque closures), so the compression step can decompose their structure. Tasks unsolved after all strategies pass to the type-directed enumerator with a cost budget of 5 and a 5-second per-task time limit.

Abstraction sleep. The corpus of solved-task programs is compressed by the Stitch anti-unifier [3]. Stitch runs pairwise anti-unification over program ASTs, replacing differing subtrees with typed holes. The MDL objective accepts an abstraction if it reduces total description length:

$$\Delta L = \text{size}(\text{pattern}) \times (\text{count} - 1) - \text{count} - n_{\text{holes}} > 0$$

Accepted abstractions are registered as callable primitives with concrete type signatures (so the enumerator can treat them as any other primitive in the next cycle) and deduplication prevents re-invention of existing patterns.

System-provenance tracking. Every abstraction tracks `systems_composed`: the union of Spelke-system tags from all primitives in its body, *plus* the systems of primitives that were generalized into holes (since holes represent the caller's contribution). This is the key fix required to detect the Carey signature: a hole-bearing abstraction that abstracts away a FORMS operation over an OBJECTS pipeline reports `systems_composed = {FORMS, OBJECTS}` even though the FORMS primitive does not literally appear in the abstract body.

Dreaming sleep. A DreamGenerator constructs synthetic programs by three modes: (30%) pipeline dreams that string OBJECTS primitives; (30%) cross-system dreams that deliberately compose OBJECTS with FORMS; and (40%) standard PCFG sampling. These programs train a two-layer NumPy MLP recognition network ($50 \rightarrow 128 \rightarrow n_{\text{prims}}$ with He initialization, ReLU, sigmoid output, binary cross-entropy loss) to predict which primitives appear in a solution given the task's input-output pairs encoded as a 50-dimensional feature vector.

3.3 The Baselines

We compare against two baselines. The **Generic DSL** (143 primitives) uses the same architecture, strategies, enumeration budget, and Stitch compressor but with flat, non-modular primitive organization: 143 generic list/grid operations without Spelke system provenance tags. This is matched in cardinality to Spelke. The **VIMRL DSL** (45 primitives, Ainooson et al. [1]) is an objectness-only DSL covering segmentation and spatial operations but without FORMS geometry or NUMBER arithmetic modules. We use VIMRL as it represents the closest prior cognitive-AI approach to ARC-AGI program synthesis.

4. Experiments

4.1 Task Sets

We evaluate on three task sets (420 tasks total):

ARC-AGI training set (400 tasks). The standard benchmark [6]. Each task has 3–6 training input-output pairs and 1–3 test pairs.

Synthetic cross-system tasks (15 tasks). We generate ARC-format tasks requiring programs of the form `(T (render_object (obj_largest (extract_objects input))))` where $T \in \{\text{rotate90, rotate180, rotate270, flip_h, flip_v}\}$, 3 instances each. Tasks use L-, T-, J-, S-, C-shaped objects with distinct coloring so that (a) the correct output is not achievable by any single FORMS primitive applied to the full input grid, and (b) no existing AST-solver strategy other than the cross-system strategy `_try_form_on_extracted_object` can solve them. These tasks are included in both cycles to allow the system to discover the cross-system abstraction.

Compounding tasks (5 tasks). Tasks of the form `(extract_tile (render_object (obj_largest (extract_objects input))))` where the input contains a solid-color rectangular object (so `extract_tile` returns its 1×1 color tile) and a single-pixel marker of a different color. These are *withheld from cycle 0*: they are presented only in cycle 1, after the OBJECTS-pipeline abstraction $\text{abs}_{0,1}$ has been discovered and registered. No cycle-0 strategy reaches `extract_tile` composed with the OBJECTS pipeline.

4.2 Claim 1: Cross-System Abstraction Discovery

We run the system for one cycle on the 15 synthetic cross-system tasks plus 400 ARC tasks (415 total). The AST solver produces 16 programs with `systems_composed  $\supset$  {FORMS, OBJECTS}`. Stitch anti-unifies these programs and discovers five abstractions; Table 2 shows the full list. Note: abstraction indices ($\text{abs}_{0,k}$) are local to each experiment run; the same index may refer to different abstractions in Table 2 (synthetic+ARC run) versus §4.5 (canonical ARC-only run).

Name	Systems	ΔMDL	Body (abbreviated)
$\text{abs}_{0,0}$	FORMS+OBJECTS	124	<code>(λinput. (•₁ (render_object (obj_largest (extract_objects input))))</code>

$\text{abs}_{0,1}$	OBJECTS	95	<code>(render_object (obj_largest (extract_objects input)))</code>
$\text{abs}_{0,2}$	OBJECTS	63	<code>(obj_largest (extract_objects input))</code>
$\text{abs}_{0,3}$	OBJECTS	43	<code>(extract_objects input)</code>
$\text{abs}_{0,4}$	FORMS	37	<code>(λinput. ((vstack (hstack input (flip_h input))) ...))</code>

Table 2. All five abstractions discovered in cycle 0. •₁ denotes a typed hole accepting any grid→grid function. $\text{abs}_{0,0}$ is the Carey-signature abstraction: FORMS+OBJECTS, MDL savings 124. $\text{abs}_{0,1}$ names the full OBJECTS extraction pipeline. The remaining three are subsystem-specific subpatterns.

Result 1. The system discovers $\text{abs}_{0,0}$ with `systems_composed = {FORMS, OBJECTS}` and MDL savings 124. This abstraction was not given to the system; it was recovered by compressing 16 cross-system programs sharing the OBJECTS-extraction backbone. The no-OBJECTS baseline discovers 0 cross-system abstractions.

4.3 Claim 2: Compounding

To demonstrate that library growth enables previously-unsolvable tasks, we constructed the *Spelke Bootstrap Suite v2* benchmark: 875 training tasks comprising 800 standard cross-system tasks (Types 1–8) plus 75 *Tier 3* compounding tasks designed to be unsolvable without the abstractions discovered in cycle 0.

Tier 3 task design. Each Tier 3 task applies one additional geometric transform (rotate90) to the output of a previously-discovered cross-system abstraction:

```

T3-A (25 tasks): rotate90(abs_0_3(input))
  abs_0_3 = render_count_colored(count_objects(...), o
bj_color(obj_largest(...)))
  Without abs_0_3: depth-5 program, unreachable in cyc
le 0

T3-B (25 tasks): rotate90(abs_0_4(input))
  abs_0_4 = tile_n(render_object(obj_largest(...)), co
unt_objects(...))
  Without abs_0_4: depth-5 program, unreachable in cyc
le 0

T3-C (25 tasks): rotate90(abs_0_5(input))
  abs_0_5 = flip_h(render_object(obj_largest(...)))
  Without abs_0_5: double-transform on extracted objec
t, no single-pass strategy

```

All three abstraction types (NUMBER+OBJECTS and FORMS+OBJECTS) are discovered in cycle 0 from the standard Type 1–8 tasks. In cycle 1, `_try_library_chains` Phase 2 composes each abstraction with `rotate90`, yielding a depth-2 program that could not be reached as depth 5+ in cycle 0.

Cycle	Train rate	Test rate	Library size	New Abs	Cross-sys Abs
0	800/875 (91.4%)	200/200 (100.0%)	206	15	12
1	875/875 (100.0%)	—	236	15	24
2	875/875 (100.0%)	—	242	3	27
3–4	875/875 (100.0%)	—	242	0	27

Table 3. Benchmark v2 compounding results (5 cycles, 875 train tasks). Cycle 0 solves 800/875; the 75 Tier 3 tasks are unsolvable without the library. Cycle 1 solves all 875: the 75 Tier 3 tasks become solvable via `_try_library_chains` composing discovered abstractions (`abs_0_3`, `abs_0_4`, `abs_0_5`) with `rotate90`. Library grows 206→236→242. The 200 held-out test tasks (Types 1–8, no Tier 3) are solved 200/200 by the base AST strategies, consistent with train-set coverage of those types. Data from `experiments/outputs/benchmark_v2_run/spelke/cycle_results.json`.

Result 2. 75 tasks unsolvable in cycle 0 become solvable in cycle 1 using abstractions discovered in cycle 0 (`abs_0_3/4/5`: NUMBER+OBJECTS and FORMS+OBJECTS cross-system patterns). Solve rate rises 91.4%→100.0%. Library grows 206→236→242, accumulating 27 cross-system abstractions over 5 cycles. This is the compounding signal: abstractions discovered in cycle 0 become first-class primitives in cycle 1, enabling tasks that were structurally inaccessible in cycle 0. The 200 held-out test tasks (Types 1–8) confirm base-system coverage; a held-out Tier 3 test set (requiring the discovered abstractions) is identified as future work.




Figure 2. Compounding mechanism. Left: the depth-4 program is unreachable within the enumerator’s cost budget in cycle 0. Right: after library learning discovers `abs0,1` (the OBJECTS extraction pipeline), the equivalent program has depth 2 and is found in milliseconds. Color coding: blue = OBJECTS, green = FORMS.

4.4 Claim 3: Spelke > Generic > VIMRL (3-Way Comparison)

We evaluate three DSLs on the 400 ARC-AGI training tasks using 3 wake-sleep cycles (enum-cost=5, enum-budget=20s, 3 cycles): the full Spelke-initialized system (143 primitives across OBJECTS, FORMS, NUMBER, and GLUE modules), a Generic matched-cardinality baseline (143 generic list/grid primitives without Spelke module structure), and the VIMRL objectness-only DSL (45 primitives, Ainooson et al. 2023 [1]) which captures object segmentation without FORMS or NUMBER modules.

System	Primitives	ARC-AGI (400)	Solve Rate	Abstractions	Cross sys Abs
Spelke (ours)	143	46 / 400	11.5%	14	1

Generic	143	42 / 400	10.5%	0	0
VIMRL [1]	45	42 / 400	10.5%	0	0
Δ (Spelke–Generic)	—	+4	+1.0pp	+14	—
Δ (Spelke–VIMRL)	—	+4	+1.0pp	+14	—

Table 4. Three-way comparison on 400 ARC-AGI training tasks (3 cycles, enum-cost=5, enum-budget=20s/task). Spelke (11.5%) outperforms Generic (10.5%) and VIMRL (10.5%) by 1.0pp. Data from [experiments/outputs/canonical_run_v2/](#).

Statistical Analysis

McNemar’s test yields $\chi^2=1.3333$, $p=0.2482$ for Spelke vs. Generic across 5 seeds ($N=400$ tasks per seed). The +4 task advantage (46 vs. 42) is consistent across all 5 seeds but does not reach conventional significance ($p<0.05$) at this scale. We interpret this as a power limitation: detecting a ~1% solve-rate difference requires approximately 1500+ tasks under McNemar’s test. Sample efficiency curves (Table 7) provide complementary monotonic evidence for the advantage.




Figure 3. Three-way comparison on 400 ARC-AGI tasks. Left: solve rate (%). Right: number of unique abstractions discovered across 3 wake-sleep cycles. Spelke discovers 14 abstractions (including 1 cross-system); both baselines discover 0.

The Generic baseline discovers zero abstractions across all three cycles: without Spelke-system provenance tracking, the compressor finds no recurring multi-primitive patterns worth compressing. The VIMRL baseline, limited to 45 objectness primitives, similarly discovers zero abstractions. This confirms that FORMS/NUMBER module structure is critical for library learning, not just task solving.

4.5 ARC-Native Cross-System Discovery and Compounding

Beyond the synthetic demonstration (§4.2–4.3), the canonical 400-task ARC run produces both cross-system discovery and compounding on *real* ARC tasks.

Cross-system abstraction on ARC. In cycle 0, `abs0,2` emerges with `systems_composed = {FORMS, OBJECTS}` and MDL savings of 25, reused across 4 programs. This is a parameterized abstraction discovered without access to synthetic tasks; the system found the FORMS+OBJECTS pattern in naturally-occurring ARC task solutions. MDL savings are lower than in the synthetic setting (25 vs. 124) because fewer ARC tasks exhibit cross-system structure, consistent with ARC’s single-concept design [6].

Compounding on ARC. Cycle 0 discovers `abs0,10`, an OBJECTS-system abstraction (`grid → grid`, MDL savings = 5). In cycle 1, the `_try_library_chains` strategy composes this abstraction with OBJECTS gravity operations to solve task `b9b7f026` via the program (`abs0,10 (gravity_right (gravity_down input))`). This task was unsolvable in cycle 0: no single strategy or enumeration path reaches the required depth-3 composition. With `abs0,10` as a cost-1 primitive, the depth-2 program is found immediately.

Cycle	Task	Reachable	Program depth	AST size	Method
0	<code>b9b7f026</code>	No	—	—	unreachable (cost budget exceeded)
1	<code>b9b7f026</code>	Yes	2	6.0	<code>abs_{0,10}</code> (OBJECTS pipeline)

Table 5. Compounding on task `b9b7f026`. `abs0,10` reduces the program from unreachable depth (≥ 3 without abstraction) to depth 2, AST size 6.0 nodes.

Cycle	Solved	New Abstractions	Cross-System	New Tasks via Compounding
0	45 / 400	14	1	—
1	46 / 400	0 (dedup)	—	+1 (b9b7f026)
2	46 / 400	0 (dedup)	—	0

Table 6. Cycle-by-cycle progression on 400 ARC-AGI tasks (canonical_run_v2). Cycle 1 solves one new task via the abstraction discovered in cycle 0. Deduplication correctly prevents re-invention of existing abstractions in later cycles.

Result 3. On 400 real ARC-AGI tasks, the Spelke system discovers 1 cross-system abstraction (FORMS+OBJECTS, $abs_{0,2}$) and achieves compounding: cycle 1 solves 1 task (b9b7f026) unsolvable in cycle 0, via $abs_{0,10}$ (OBJECTS abstraction) reducing an unreachable depth-3 program to a depth-2 program. Library grows from 143 to 171 primitives; Generic and VIMRL libraries remain static. Data: experiments/outputs/canonical_run_v2/.




Figure 4. Left: library size across wake-sleep cycles. Spelke grows from 143 to 171 primitives; Generic and VIMRL remain static. Right: composition of the 14 discovered abstractions by Spelke system provenance.

4.6 Sample Efficiency

We measure solve rate as a function of the number of training tasks (task_count $\in \{10, 25, 50, 100, 200, 400\}$, 3 seeds each)

to characterize how quickly the Spelke advantage emerges with sample size. Data from experiments/outputs/sample_efficiency_full_20260526_205111/sample_efficiency_results.json.

Tasks (N)	Spelke	Generic	Δ
10	10.0%	10.0%	0.0%
25	10.7%	10.7%	0.0%
50	11.3%	10.0%	+1.3%
100	12.0%	10.3%	+1.7%
200	11.0%	9.8%	+1.2%
400	11.2%	10.5%	+0.7%

Table 7. Sample efficiency: mean solve rate across 3 seeds. Spelke maintains advantage at all task counts ≥ 50 . Rates from sample_efficiency_results.json.




Figure 5. Sample efficiency: solve rate as a function of training set size (mean across 3 seeds). The Spelke advantage emerges at $N \geq 50$ tasks and is maintained across all larger task counts, peaking at +1.7% at $N=100$. The shaded region highlights the persistent gap. At $N \leq 25$ both systems perform identically, suggesting that cross-system structure requires a minimum corpus size to produce compressible patterns.

4.7 Extended Spelke DSL (AGENTS + PLACES)

We evaluate the full 170-primitive Spelke DSL extended with AGENTS and PLACES (28 additional primitives beyond the 143-primitive base) on 400 ARC-AGI tasks, 1 cycle. The extended DSL achieves 45/400 (11.2%), slightly below the 143-primitive base (46/400, 11.5%): adding AGENTS and PLACES does not improve ARC performance. Library grows to 198 primitives (170 base + 28 discovered abstractions). AGENTS and PLACES primitives do not fire on the ARC training distribution: zero tasks require agent-trajectory or spatial-navigation reasoning, consistent with ARC's single-concept design [6]. Data from

```
experiments/outputs/phase2_spelke/spelke/results.json.
```

4.8 Curriculum Pre-training with Cross-System Tasks

We construct a curriculum of 120 synthetic ARC-format tasks spanning four cross-system patterns: **NUMBER+OBJECTS** (30 tasks, canvas sizes 8×8 – 15×15): count N same-colored objects, render as $1\times N$ colored row using `count_objects`; **FORMS+OBJECTS+PLACES** (30 tasks, 8×8): extract largest object, apply geometric transform (rotate/flip), place in target quadrant; **NUMBER+OBJECTS replicate** (30 tasks, canvas sizes 8×8 – 15×15): extract largest object, tile it N times horizontally where N = total object count; **count_cells** (30 tasks, 3×3 – 4×4): N non-zero cells arranged in a connected region, render as $1\times N$ colored row using `count_cells`; this pattern covers ARC tasks where the count is of total pixels rather than connected components, and directly matches ARC task d631b094. Task sizes span the ARC distribution to encourage size-invariant abstractions.

Cross-system abstraction discovery. After one curriculum pre-training cycle over 140 tasks (120 cross-system + 15 synthetic cross-system + 5 compounding), the system discovers 10 cross-system abstractions (of 15 total): 8 NUMBER+OBJECTS, 1 FORMS+OBJECTS, and 1 FORMS+OBJECTS+PLACES. The top NUMBER+OBJECTS abstractions have MDL savings of 434, 434, 405, 405, and 376 bits, all with `is_cross_system = True`; the Carey signature holds at the primitive level. Notably, the `count_cells` curriculum tasks yield a cross-system abstraction using the `count_cells` NUMBER primitive:

```
(λinput. (render_count_colored (count_cells input) (obj_color (obj_largest (extract_objects input))))).
```

 Spelke solves 135/140 curriculum tasks; the Generic DSL baseline solves 0. The library grows to 193 primitives (178 base + 15 discovered abstractions). Data from `experiments/outputs/curriculum_v2_arc_run/pretrain/`.

Key cross-system abstractions:

- *abs_0_2* (savings=434):

```
(λinput. ((render_count_colored (count_objects (extract_objects input)) (obj_color (obj_largest (extract_objects input))))):
```

 counts all OBJECTS-system components and renders that NUMBER-system count as a colored row.
- *abs_0_3* (savings=434):

```
(λinput. ((tile_n (render_object (obj_largest (extract_objects input))) (count_objects (extract_objects input))))):
```

 extracts

the largest OBJECTS-system shape and tiles it by the NUMBER-system object count.

- *abs_0_6* (savings=376): `count_cells`-based NUMBER+OBJECTS bridge; renders the total pixel count (not component count) as a $1\times N$ row.

Compounding. Using cycle 1 abstractions, the system solves 5/5 compounding tasks that were unsolvable before. These tasks require first applying a cross-system abstraction and then composing the result. Abstractions from curriculum tasks transfer to solve previously-unsolved problems. Advantage over Generic DSL: +135 tasks (135 vs. 0 on curriculum).

ARC transfer (confirmed). We run a 3-cycle ARC experiment with curriculum pre-training on all 400 ARC-AGI training tasks (`enum-cost=6`, `enum-budget=10s`, Rust enumerator). Results from `experiments/outputs/curriculum_v2_arc_run/`. The system solves **47/400** ARC tasks (vs. 46/400 baseline without curriculum pre-training), confirmed by `verify_transfer.py` (5/5 checks). Critically, ARC task **d631b094**, previously unsolved, is now solved in cycle 0 using the cross-system program:

```
(λinput. (render_count_colored (count_cells input) (obj_color (obj_largest (extract_objects input))))).
```

 This program bridges NUMBER (`count_cells`, `render_count_colored`) and OBJECTS (`obj_color`, `obj_largest`, `extract_objects`) systems, exactly the Carey cross-system combination taught in the curriculum. The curriculum pre-trained library grew to 193 primitives.

4.9 PERSONS System and Full Six-System Comparison

We complete the Spelke framework by implementing the sixth and final core knowledge system: **PERSONS** [20]. PERSONS primitives encode the social-cognition signatures infants use to distinguish persons from objects: asymmetry/directionality detection (`is_agent_obj`, `is_target_obj`), agent counting (`agent_count`), intentional-action attribution (`obj_facing_dir`, `point_toward`), social-distance measurement (`social_dist`), agent/target isolation (`extract_agents`, `extract_targets`, `nearest_agent`, `nearest_target`), and theory-of-mind proxies (`goal_reachable`, `mirror_intent`). This completes the hexagon: OBJECTS, FORMS, NUMBER, AGENTS, PLACES, PERSONS.




Figure 6. The complete Spelke Core Knowledge Framework as implemented in bootstrap-substrate. Phase 1 systems (OBJECTS, FORMS, NUMBER; colored) contribute to ARC task solving; Phase 2 systems (AGENTS, PLACES, PERSONS; gray) encode domain-specific cognitive distinctions validated on curriculum tasks. The red dashed line marks the FORMS+OBJECTS cross-system abstraction, the Carey signature discovered by MDL compression. Total: 200 base primitives across 6 systems.

Six-system DSL. The full 6-system DSL (200 base primitives) achieves 47/400 (11.75%) on ARC-AGI training tasks over 2 cycles (enum_cost=4, enum_budget=2s; note: these parameters differ from the canonical comparison in §4.4, so this figure is not directly comparable to the 46/400 baseline). Library grows from 200 to 214 primitives (14 new abstractions discovered; 1 cross-system: FORMS+OBJECTS). As with AGENTS and PLACES (Section 4.7), PERSONS primitives do not fire on native ARC tasks; ARC tasks are designed around geometric/physical reasoning, not social cognition. This null result is expected and scientifically honest: the claim is that the *mechanism* for social bootstrapping exists, not that ARC tests it. Data from `experiments/outputs/six_system_run/`.

System	Primitives	Phase	ARC impact
OBJECTS	38	Phase 1	Core (46+/400)
FORMS	31	Phase 1	Core
NUMBER	41	Phase 1	Core
AGENTS	11	Phase 2	Null (0 ARC tasks)
PLACES	16	Phase 2	Null (0 ARC tasks)
PERSONS	12	Phase 2	Null (ARC not social)

PERSONS curriculum tasks. On 30 PERSONS+OBJECTS curriculum tasks, the AST solver with PERSONS enabled solves 30/30 tasks (100%); the no-PERSONS generic DSL solves 0/30. Stitch discovers 9 abstractions, of which 4 are cross-system with PERSONS in `systems_composed`. The top cross-system abstractions from this curriculum:

- `abs_0_3` (savings=153): `(λinput. ((render_count_colored (agent_count input)) (obj_color (nearest_agent input))))` — NUMBER+OBJECTS+PERSONS cross-system abstraction.
- `abs_0_4` (savings=139): `((render_count_colored (agent_count input)) (obj_color (nearest_agent input)))` — NUMBER+OBJECTS+PERSONS cross-system abstraction.

The **TYPE F** abstraction `(λinput. ((render_count_colored (agent_count input)) (obj_color (nearest_agent input))))` bridges three systems: PERSONS (`agent_count`, `nearest_agent`), NUMBER (`render_count_colored`), and OBJECTS (`obj_color`). This is the Carey signature applied to social cognition: the system learns that “count the agents and render that many colored cells” is a reusable cross-domain operation, composing PERSONS-system detection with NUMBER-system cardinality and OBJECTS-system color attribution. Data from `experiments/outputs/persons_curriculum_demo/curriculum_demo_results.json`.

Completing the framework. With all six Spelke core knowledge systems implemented and evaluated, we can claim the first computational instantiation of Spelke’s full core knowledge framework as a program-induction library. The null ARC results for AGENTS, PLACES, and PERSONS confirm ARC’s single-concept design philosophy; the positive curriculum results confirm that each system’s primitives encode genuine cognitive distinctions that produce cross-system abstractions when those distinctions are exercised.

5. Discussion

5.1 The Carey Signature

Carey’s account requires three things: (i) placeholder symbols from distinct systems; (ii) cross-system modelling processes that compose them; (iii) the resulting concept has expressive power exceeding either source. We assess each using the cross-system abstractions discovered in our experiments.

(i) *Placeholders.* The hole $\bullet_1$ in `abs0,0` (synthetic, MDL savings 124) and `abs0,2` (ARC-native, MDL savings 25) is exactly a Quinean placeholder: a typed slot (type `grid`→`grid`) that the system uses before fixing its value. The hole ranges over the full space of FORMS transforms, a kind of intentional indeterminacy fundamental to the bootstrapping account.

(ii) *Cross-system composition.* The body of both abstractions combines the OBJECTS pipeline (`extract_objects` → `obj_largest/obj_smallest` → `render_object`) with a FORMS placeholder. Neither system can express this on its own: OBJECTS lacks geometric transforms; FORMS has no notion of object individuation. The composition is the cross-system modelling process.

(iii) *Expressive power.* The compounding results (§4.3, §4.5) provide the operational definition. When OBJECTS pipeline abstractions become first-class primitives, the system reaches programs of previously inaccessible depth. On synthetic tasks (§4.1), `extract_tile(abs0,1(input))` becomes reachable: `extract_tile` is a FORMS primitive applied to an OBJECTS abstraction, a depth-2 program whose underlying OBJECTS pipeline would otherwise exceed the cost budget; this is the cross-system composition Carey predicts. On ARC (§4.5), task `b9b7f026` is solved in cycle 1 via `(abs0,10 (gravity_right (gravity_down input)))`: an OBJECTS-system composition that is unreachable as a depth-3 program in cycle 0 but becomes depth-2 once `abs0,10` is registered as a cost-1 primitive. Both cases satisfy condition (c): at least one task becomes solvable in the next cycle that was structurally inaccessible before.

5.2 Cross-System Structure in ARC

A finding with theoretical significance: of 400 ARC training tasks, only one naturally produces a program that combines FORMS and OBJECTS system primitives. ARC tasks are designed as clean, single-concept puzzles [6], and this design constraint limits cross-system opportunity. The canonical run discovers one cross-system abstraction on ARC (`abs0,2`, MDL savings 25, reuse 4), demonstrating that the mechanism works on real data, but the signal is sparse. This has implications for benchmark design: if the goal is to test PrimaLearn at scale, evaluation requires tasks with richer cross-system structure.

5.3 Limitations and Honest Assessment

Statistical power. The +4 task advantage (46 vs. 42 tasks, 11.5% vs. 10.5%) is *not statistically significant* at conventional thresholds (McNemar $\chi^2=1.3333$, $p=0.2482$). The 400-task ARC benchmark is underpowered for detecting 1% solve-rate differences. A larger task set (≥ 1500 tasks) is required to confirm this advantage at conventional significance levels.

AGENTS and PLACES primitives. The 170-primitive full Spelke DSL achieves 45/400 (11.2%) vs. 46/400 (11.5%) for 143-primitive Spelke. AGENTS and PLACES primitives do not fire on the ARC training distribution, which encodes single-concept tasks by design [6]. This is a null result and an expected limitation of the benchmark scope.

Synthetic task design. The 15 cross-system and 5 compounding tasks (§4.1) and the 875-task Spelke Bootstrap Suite v2 (§4.3) were designed by the authors to exhibit the target phenomenon. This is correct experimental methodology (one cannot test bootstrapping without tasks that require it), but limits external validity. A key open question is whether naturally-occurring visual reasoning benchmarks contain sufficient cross-system structure to replicate these results.

Enumerator speed. We built a Rust enumerator achieving 4.6–9.4× speedup over the Python baseline (§4.8). This is an improvement, but it remains well below the original DreamCoder Rust enumerator, which uses task-parallel enumeration, PCFG-guided top-down search, and bitset memoization to achieve approximately 1000× Python speed. Cross-system programs at depth ≥ 6 remain unreachable within practical budgets. A fully optimized Rust implementation with these techniques would substantially expand the discovery space. This is an engineering limitation, not a theoretical one.

Baseline fairness. The Generic baseline uses a matched-cardinality DSL (143 primitives) without Spelke module structure. The VIMRL baseline uses 45 objectness primitives following Ainooson et al. [1]. The Spelke vs. Generic solve-rate gap (+1.0pp) is modest; the more significant gap is in library growth: 14 abstractions vs. 0. The VIMRL parity with Generic on solve rate but zero abstractions confirms that FORMS/NUMBER module structure is critical for library learning, not just task solving.

Single cycle of compounding. We demonstrate one cycle of compounding. The DreamCoder account predicts that this compounds across many cycles, with each cycle's abstractions enabling the next. Demonstrating multi-cycle bootstrapping

requires a task set with hierarchical cross-system structure, which we leave to future work.

5.4 Relation to Prior Work

DreamCoder [8] demonstrates library learning on list processing, string editing, and scene understanding but does not ground primitives in Spelke systems and does not measure cross-system provenance. Stitch [3] provides the efficient anti-unification backbone we build on. VIMRL [1] is the closest cognitive-AI predecessor but lacks FORMS/NUMBER systems and library learning. Chollet's ARC paper [6] explicitly calls for systems grounded in core knowledge; this work is a direct response to that call.

6. Conclusion

We have presented the first computational instantiation of Carey's Quinian bootstrapping theory within a neural program-synthesis system. By initializing the DSL with Spelke's core-knowledge systems, running a DreamCoder wake-sleep loop with Stitch compression, and tracking cross-system provenance through every abstraction, we demonstrate that:

1. The system spontaneously discovers `abs0,2`, a FORMS+OBJECTS combinator with MDL savings 25 on real ARC tasks, and `abs0,0` with savings 124 on synthetic cross-system tasks, both matching the Carey signature of a placeholder spanning incommensurable systems.
2. An OBJECTS abstraction (`abs0,10`) discovered in cycle 0 enables solving task `b9b7f026` in cycle 1 by composing with OBJECTS gravity operations, a depth-3 program unreachable in cycle 0 that becomes depth-2 once the abstraction is registered as a cost-1 primitive. On the Spelke Bootstrap Suite v2

(§4.3), 75 previously unsolvable Tier 3 tasks become solvable in cycle 1 via three discovered cross-system abstractions.

3. In a three-way comparison on 400 ARC-AGI tasks (5 seeds), the Spelke-initialized system achieves 46/400 (11.5%) versus 42/400 (10.5%) for both the matched-cardinality Generic DSL and the VIMRL objectness-only DSL. McNemar's test: $\chi^2=1.3333$, $p=0.2482$. Spelke discovers 14 unique library abstractions across 3 cycles (library 143→171); both baselines discover 0. Data: `experiments/outputs/canonical_run_v2/`.

PrimaLearn is no longer purely a verbal theory. It is a measurable property of a running system, with an operational definition: an abstraction bootstraps if and only if it (a) combines primitives from at least two distinct Spelke systems, (b) was not given to the system but was recovered by MDL compression, and (c) enables solving at least one task in the next cycle that was unsolvable in the prior cycle. All three conditions are met here.

Future directions include: multi-cycle bootstrapping on hierarchically structured benchmarks (Tier 4 tasks requiring Tier 3 abstractions to build a true staircase); optimizing the existing Rust enumerator with task-parallel search and PCFG-guided enumeration to reach programs at depth ≥ 6 ; scaling to ≥ 1500 tasks for statistical significance; and applying the provenance-tracking machinery to large language models whose implicit knowledge systems remain largely unmapped.

Acknowledgments

The author thanks the open-source communities behind the ARC-AGI benchmark and the Stitch anti-unification library. This work was conducted independently. Code and data will be made available at a public repository upon publication.

References

- [1] Ainooson, J., Sanyal, D., Michelson, J. P., Yang, Y., & Kunda, M. (2023). An approach for solving tasks on the abstract and reasoning corpus. *AAAI Workshop on Representation Learning for Responsible Human-Centric AI*.
- [2] Bates, C. J., Yildirim, I., Tenenbaum, J. B., & Battaglia, P. W. (2019). Modeling human intuitions about liquid flow with particle-based simulation. *PLoS Computational Biology*, 15(7).
- [3] Bowers, M., Olausson, T. X., Wong, L., Grand, G., Tenenbaum, J. B., Ellis, K., & Solar-Lezama, A. (2023). Top-down synthesis for library learning. *Proceedings of the ACM on Programming Languages*, 7(POPL), 1–29.
- [4] Carey, S. (2009). *The Origin of Concepts*. Oxford University Press.
- [5] Carey, S., & Spelke, E. S. (1994). Domain-specific knowledge and conceptual change. In L. Hirschfeld & S. Gelman (Eds.), *Mapping the Mind* (pp. 169–200). Cambridge University Press.
- [6] Chollet, F. (2019). On the measure of intelligence. *arXiv preprint arXiv:1911.01547*.
- [7] Ellis, K., Morales, L., Sablé-Meyer, M., Solar-Lezama, A., & Tenenbaum, J. B. (2018). Learning libraries of subroutines for neurally-guided Bayesian program induction. *Advances in Neural Information Processing Systems*, 31.
- [8] Ellis, K., Wong, C., Nye, M., Sablé-Meyer, M., Morales, L., Hewitt, L., Cary, L., Solar-Lezama, A., & Tenenbaum, J. B. (2021). DreamCoder: Bootstrapping inductive program synthesis with wake-sleep library learning. *Proceedings of the 42nd ACM SIGPLAN Conference on Programming Language Design and Implementation*, 835–850.
- [9] Feser, J. K., Chaudhuri, S., & Dillig, I. (2015). Synthesizing data structure transformations from input-output examples. *ACM SIGPLAN Notices*, 50(6), 229–239.

- [10] Gelman, R., & Gallistel, C. R. (1978). *The Child's Understanding of Number*. Harvard University Press.
- [11] Hodel, M. (2024). Addressing the abstraction and reasoning corpus via procedural example generation. *arXiv preprint arXiv:2404.07353*.
- [12] Johnson, J., Hariharan, B., van der Maaten, L., Fei-Fei, L., Lawrence Zitnick, C., & Girshick, R. (2017). CLEVR: A diagnostic dataset for compositional language and elementary visual reasoning. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2901–2910.
- [13] Lake, B. M., Salakhutdinov, R., & Tenenbaum, J. B. (2015). Human-level concept learning through probabilistic program induction. *Science*, 350(6266), 1332–1338.
- [14] Lake, B. M., Ullman, T. D., Tenenbaum, J. B., & Gershman, S. J. (2017). Building machines that learn and think like people. *Behavioral and Brain Sciences*, 40.
- [15] Liang, P., Jordan, M. I., & Klein, D. (2010). Learning programs: A hierarchical Bayesian approach. *Proceedings of the 27th International Conference on Machine Learning*, 639–646.
- [16] Mirchandani, S., Karamcheti, S., & Sadigh, D. (2023). Large language models as general pattern machines. *arXiv preprint arXiv:2307.04721*.
- [17] Muggleton, S. (1991). Inductive logic programming. *New Generation Computing*, 8(4), 295–318.
- [18] Overlan, C. M., Jacobs, R. A., & Piantadosi, S. T. (2017). Learning abstract visual concepts via probabilistic program induction in a Language of Thought model. *Cognition*, 168, 320–334.
- [19] Piantadosi, S. T., Tenenbaum, J. B., & Goodman, N. D. (2012). Bootstrapping in a language of thought: A formal model of numerical concept learning. *Cognition*, 123(2), 199–217.
- [20] Spelke, E. S., & Kinzler, K. D. (2007). Core knowledge. *Developmental Science*, 10(1), 89–96.
- [21] Xu, K., Li, J., Zhang, M., Du, S. S., Kawakami, K., Bengio, Y., & Courville, A. (2024). LLMs and the abstraction and reasoning corpus: Successes, failures, and the importance of object-based representations. *arXiv preprint arXiv:2305.18354v3*.